

GETTING STARTED ON THE VACC

Vermont Advanced Computing Center

—
Shelly Johnson
—
Senior Research Computing Facilitator

—
Bennet Fauber
—
Senior Research Computing Systems Engineer

High Performance Computing

- A High Performance Computing (HPC) cluster is a collection of many computers networked together
- Each computer is called a node.
- HPC clusters run computations managed by a scheduler
- The scheduler keeps track of available resources, allowing requests to be efficiently assigned to compute resources
 - memory, processing power (cpu or gpu), time, or some combination

Node Types

Login Node

When you log onto a cluster you land on a login node. Edit scripts, run short tests, and

```
[ <netID>@vacc-user1 ~] $
```


Rough guidelines: the test will run < 5 minutes, use less than 8-12 GB of memory, and u
(less is better)

If more resources are needed, use a regular job or an interactive job or Open OnDemand

Compute Node

Compute nodes are where all the heavy lifting is done. Once on a login node, you can s
jobs either in batch form or interactively to the compute nodes.

The VACC Cluster



Slurm

All jobs are run on the compute nodes using the Slurm Workload Manager. Slurm needs to have the following information about what you want to run:

- The partition (group of nodes) to use
- Number of CPUs
- Amount of memory
- Amount of time needed to run
- Whether you need a GPU, and if so, how many

These are the boxes you will fill in when starting a job

For further information about Slurm, see the User Guide

<https://slurm.schedmd.com/quickstart.html>

Modules on VACC

- Modules are used to modify your environment to access software packages installed on the cluster
- Modules can be loaded on the command line or in your batch script
 - It is preferred to load them in your batch script to have a record of which module loaded if there is a problem with your job

Load Software on VACC: MODULES

Show which modules are **available**:

```
$ module av
```

Load a module and list the loaded modules (including dependencies)

```
$ module load Rtidyverse/4.3.2
```

```
$ module load rstudio
```

```
$ module list
```

Unload modules:

```
$ module unload Rtidyverse rstudio
```

```
$ module purge
```

Load Software on VACC: Spack

Show which packages are available:

```
$ spack find
```

Load a package and list the loaded software (including dependencies)

```
$ spack load bowtie@1.2.3
```

```
$ spack find --loaded
```

Unload a Spack package:

```
$ spack unload bowtie
```


Open OnDemand (OOD)

- Submit jobs via a web interface: <https://vacc-ondemand.uvm.edu>
- Choose between a Desktop or any of the following apps:
 - Jupyter Notebook
 - Rstudio
 - MATLAB* (coming soon)

Guide to using the VACC Open OnDemand

<https://www.uvm.edu/vacc/kb/knowledge-base/ondemand/>

Creating a Batch Script for Slurm

- A batch script is the script used to submit a job to the cluster. Batch scripts end with either *.sbat* or *.sh*
- The example shown is a simple batch script that loads a Python module and runs a Python script.
- Modules should be loaded after the `#SBATCH` lines and before any commands

```
#!/bin/bash
#SBATCH --job-name=example
#SBATCH --time=00:15:00
#SBATCH --ntasks=2
#SBATCH --mem=8G
#SBATCH --nodes=1
#SBATCH --account=<account>
#SBATCH --partition=bluematter
#SBATCH --mail-user=<netid>
#SBATCH --mail-type=BEGIN,END

module purge
module load python3.11-anaconda
module list

python my_python_script.py
```

Submitting a Job to the Cluster

Check your that your script specifies all required resources, the account partition are accurate, and all of your modules loaded

To submit your job to the cluster

```
$ sbatch <script_name>.sbat (or .sh)
```

Check the status of your job

```
$ squeue -u <user>
```

JOBID	PARTITION	NAME	USER	ST	TIME	NODES	NODELIST (REASON)
325511	bluemoon	testjob	mjohns89	R	0:00	1	node200

Helpful Commands

```
my_accounts
```

```
: list your cluster accounts
```

```
my_job_statistics <jobID>
```

```
: show resources used for a comp
```

```
squeue -u <user>
```

```
: check all jobs for a specific
```

```
scancel <jobID>
```

```
: cancel a job
```

Resources

VACC Knowledge Base:

<https://vacckb.helpline.w3.uvm.edu/vacc/kb/>

Open OnDemand (web interface):

<https://vacc-ondemand.uvm.edu/>

Open a support ticket: vacchelp@uvm.edu